

Introduction to Reinforcement Learning



Na Li



February 28, 2020

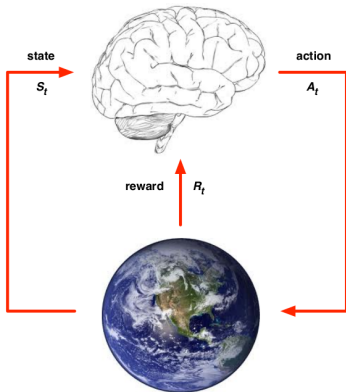
Outline

- ▶ About Reinforcement Learning
- ▶ An Example
- ▶ Limitations
- ▶ Categories
- ▶ Discrimination: Unsupervised, Supervised vs. RL
- ▶ Applications

1. About Reinforcement Learning



Introduction



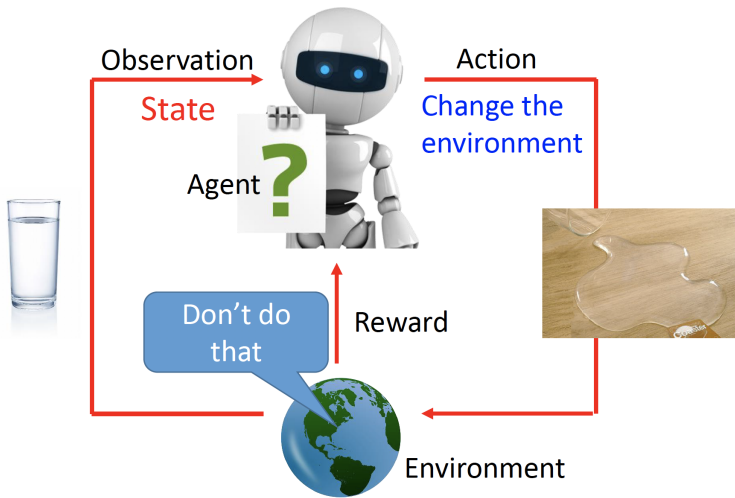
► Agent:

- Executes action A_t .
- Receives observation S_t .
- Receives reward R_t .

► Environment:

- Receives action A_t .
- Emits observation S_{t+1} .
- Emits reward R_{t+1} .

Scenario of Reinforcement Learning



Scenario of Reinforcement Learning



Core idea

Core idea of RL:

- ▶ Interacts with the environment.
- ▶ Learns from experience.
- ▶ The target is to get the maximum expected cumulative rewards.

Reward hypothesis

Expected cumulative rewards in state S in time t :

$$V(S) = E_{\pi}[G_t|S]$$

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots (\gamma \in [0, 1))$$

- ▶ the greater the γ , the long-term rewards are more concerned.
- ▶ MC method: using random samples to estimate expectations. \rightarrow TD: for efficiency.

Elements

policy $\pi_{\theta}(A|S)$

- ▶ A policy is the agent's behavior.
- ▶ It is a map from state to action.
- ▶ E.g. a function, lookup table.

Elements

reward signal $R(S, A)$

- ▶ Short-term.
- ▶ The primary basis for altering the policy.

Elements

value function $V(S)$

- ▶ A prediction of future reward, which is long-term rewards.
- ▶ Used to evaluate the goodness/badness of states.
- ▶ And therefore to select between actions.

Elements

model(optional)

- ▶ A model predicts what the environment will do next.
- ▶ \mathcal{P} predicts the next state, i.e.
 $\mathcal{P}_{ss'}^a = P[S'|S, A]$
- ▶ \mathcal{R} predicts the next reward, i.e.
 $\mathcal{R}_s^a = P[R'|S, A]$

2. An Example



An Example






Figure: House Structure


An Example



Algorithm

1. Set the λ and reward matrix R .
2. Initialize Q as zero.
3. For each episode:
 - (a) select a random initial state S .
 - (b) Do While the goal state hasn't been reached.
 - (i) Select one possible action A from current state.
 - (ii) Using A , considering going to the next state S'
 - (iii) Considering the next action A' .
 - (iv) Compute:
$$Q(S, A) = R(S, A) + \lambda \max_{A'} Q(S', A')$$
 - (v) Set the next state as the current state.

Initialization




$$R = \begin{bmatrix} -1 & -1 & -1 & -1 & 0 & -1 \\ -1 & -1 & -1 & 0 & -1 & 100 \\ -1 & -1 & -1 & 0 & -1 & -1 \\ -1 & 0 & 0 & -1 & 0 & -1 \\ 0 & -1 & -1 & 0 & -1 & 100 \\ -1 & 0 & -1 & -1 & 0 & 100 \end{bmatrix}, Q = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

Episode 1:

- (a) Select a random initial state $S = 1$.
- (b) Select one possible action $A = 5$.
- (c) Next state $S' = 5$; Next action $A' = 1, 4, 5$.
- (d) Compute:
 $Q(1, 5) = R(1, 5) + 0.8 * \max\{Q(5, 1), Q(5, 4), Q(5, 5)\} = 100 + 0.8 * \max(0, 0, 0) = 100$.

(e) $Q = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 100 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$




- (f) $S = 5$, state '5' is the goal state and end.

Episode 2:

- (a) Select a random initial state $S = 3$.
- (b) Select one possible action $A = 1$.
- (c) Next state $S' = 1$; Next action $A' = 3, 5$.
- (d) Compute:

$$Q(3, 1) = R(3, 1) + 0.8 * \max\{Q(1, 3), Q(1, 5)\} = 0 + 0.8 * \max(0, 100) = 80.$$

(e) $Q = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 100 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 80 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$




- (f) $S = 1$, state '1' is not the goal state and continue.

Episode 2:

- (a) Current state $S = 1$.
- (b) Select one possible action $A = 5$.
- (c) Next state $S' = 5$; Next action $A' = 1, 4, 5$.
- (d) Compute:
 $Q(1, 5) = R(1, 5) + 0.8 * \max\{Q(5, 1), Q(5, 4), Q(5, 5)\} = 100 + 0.8 * \max(0, 0, 0) = 100$.

(e) $Q = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 100 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 80 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$



- (f) state '5' is the goal state and end.


Episode N:

Convergence result: $Q = \begin{bmatrix} 0 & 0 & 0 & 0 & 400 & 0 \\ 0 & 0 & 0 & 320 & 0 & 100 \\ 0 & 0 & 0 & 320 & 0 & 0 \\ 0 & 400 & 256 & 0 & 400 & 0 \\ 320 & 0 & 0 & 320 & 0 & 500 \\ 0 & 400 & 0 & 0 & 400 & 500 \end{bmatrix}$

Normalization: $Q = \begin{bmatrix} 0 & 0 & 0 & 0 & 0.8 & 0 \\ 0 & 0 & 0 & 0.64 & 0 & 1 \\ 0 & 0 & 0 & 0.64 & 0 & 0 \\ 0 & 0.8 & 0.51 & 0 & 0.8 & 0 \\ 0.64 & 0 & 0 & 0.64 & 0 & 1 \\ 0 & 0.8 & 0 & 0 & 0.8 & 1 \end{bmatrix}$

Result:

$$Q = \begin{bmatrix} 0 & 0 & 0 & 0 & 0.8 & 0 \\ 0 & 0 & 0 & 0.64 & 0 & 1 \\ 0 & 0 & 0 & 0.64 & 0 & 0 \\ 0 & 0.8 & 0.51 & 0 & 0.8 & 0 \\ 0.64 & 0 & 0 & 0.64 & 0 & 1 \\ 0 & 0.8 & 0 & 0 & 0.8 & 1 \end{bmatrix}$$



3. Limitations



Limitations

Reward delay.

- ▶ Playing Games.

- Only “fire” obtains reward.
- Although the moving before “fire” is important.

- ▶ Playing Go.

- Gaining reward only after winning.
- Every step is very important.

Limitations

Exploration & Exploitation.


- ▶ Exploration finds more information about the environment.
- ▶ Exploitation exploits known information to maximise reward.
- ▶ It is usually important to explore as well as exploit.
- ▶ Restaurant Selection:
Exploitation: Go to your favourite restaurant.
Exploration: Try a new restaurant.
- ▶ Online Advertisements:
Exploitation: Show the most successful advert.
Exploration: Show a different advert.

4. Categories



Model-free vs. Model-based

- ▶ Model-free: learns the optimal policy.
e.g. Q-learning, Sarsa, Policy-gradient, DQN, A3C, ...
- ▶ Model-based: models the environment, then based on that model, chooses the most appropriate policy.



Model-free vs. Model-based

► Model-free:

- Gets the next state by interacting with the environment.
- Learns by trial and error.

► Model-based:

- Gets transition probability and return function directly.
- Given current state S and action A , we can know the next state S' and the value V directly.

Value-based vs. Policy-based

► Value-based:

- Learns the values of actions (strategy evaluation).
- Then selects actions based on their estimated action values (strategy optimization).
- E.g., Q-learning, Sarsa, DQN.

► Policy-based:

- Learns a parameterized policy that can select actions without consulting a value function.
- E.g., Policy-gradient, A3C, DDPG.
- *A value function may still be used to learn the policy parameter, but is not required for action selection.*

Value-based vs. Policy-based

Value-based: obtains $Q_{\pi}(S, A)$ firstly, then policy is:

$$\pi_* = \arg \max_{\pi} Q_{\pi}(S, A)$$

Policy-based: optimizes for the long term reward directly:

$$\begin{aligned} J(\theta) &= \sum_{S \in \mathcal{S}} d_{\pi(\theta)}(S) V_{\pi(\theta)}(S) \\ &= \sum_{S \in \mathcal{S}} (d_{\pi(\theta)} \sum_{A \in \mathcal{A}} \pi(A|S, \theta) Q_{\pi}(S, A)) \end{aligned}$$

Value-based vs. Policy-based

- ▶ Value-based: for low-dimensional, discrete action space cases.
- ▶ Policy-based: for continuous action space cases.

Off-policy vs. On-policy

- ▶ Behavior strategy: guides individuals to produce actual interaction with the environment.
- ▶ Target strategy: evaluates the state or action value.

Off-policy vs. On-policy

- ▶ Off-policy: Behavior strategy and target strategy is different, e.g. Q-learning, DQN.

$$\begin{aligned}Q(S, A) &\leftarrow Q(S, A) + \alpha[R(S, A) + \gamma \max_a Q(S', a) - Q(S, A)] \\&= (1 - \alpha)Q(S, A) + \alpha[R(S, A) + \gamma \max_a Q(S', a)]\end{aligned}$$

- ▶ On-policy: Behavior strategy and target strategy is same, e.g., Sarsa:

$$\begin{aligned}Q(S, A) &\leftarrow Q(S, A) + \alpha[R(S, A) + \gamma Q(S', A') - Q(S, A)] \\&= (1 - \alpha)Q(S, A) + \alpha[R(S, A) + \gamma Q(S', A')]\end{aligned}$$

Off-policy vs. On-policy

► Off-policy:

- A daring strategy.
- Chooses the direction of maximizing Q when updating Q .
- Then chooses the action again in the next state.

► On-policy:

- A conservative strategy.
- Planning actions for the future when updating the Q value.

5. Discrimination: Unsupervised, Supervised vs. RL




Discrimination

- ▶ Unsupervised: finds hidden structures.
- ▶ Supervised: finds mapping rules – learns from teacher(label).
- ▶ RL: finds feedback mechanisms – learns from experience.

Discrimination

- Supervised: **Learns from teacher.**



Next Move: “3-3”



Next Move: “2-2”

Discrimination

- ▶ RL: **Learns from experience.**

First move \Rightarrow ... many moves ... \Rightarrow Win!

First move \Rightarrow ... many moves ... \Rightarrow
Lose!

6. Applications



Toolkit & Demo

- ▶ Toolkit:
Gym: <http://gym.openai.com/>
Universe:
<https://openai.com/blog/universe/>
- ▶ Demo:
Cartpole & Flappybird.

Q & A